

CS 499/599: Machine Learning Security

02.07: Data Poisoning

Mon/Wed 12:00 – 1:50 pm

Sanghyun Hong

sanghyun.hong@oregonstate.edu



Oregon State
University

SAIL
Secure AI Systems Lab

Notice

- Due dates
 - Written Paper Critiques (on the 9th)
- Sign-up (on Canvas)
 - Scribe Lecture Note
 - In-class Paper Presentation / Discussion

Topics for Today

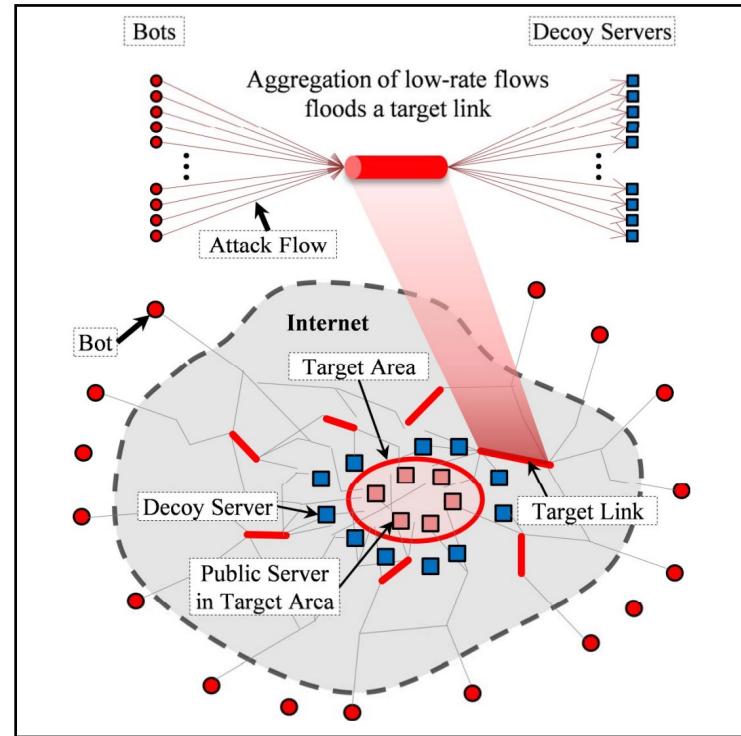
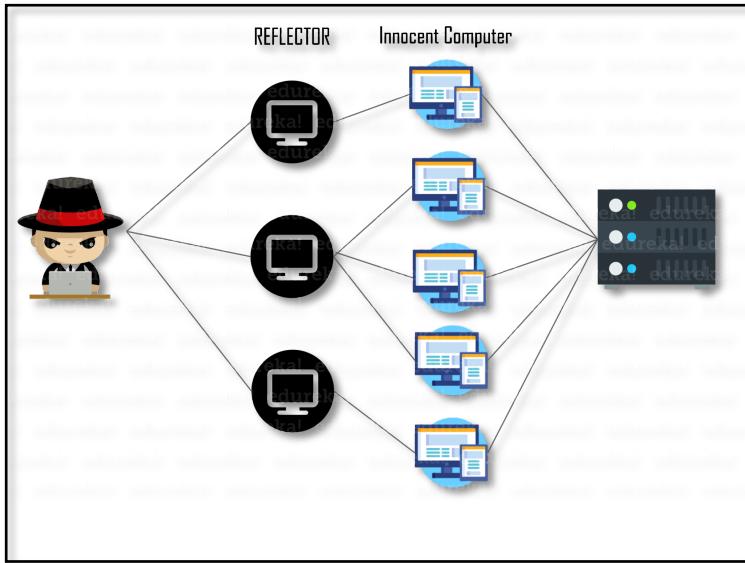
- Data Poisoning
 - Exploitations
 - Spam filtering
 - DDoS detection
 - Conclusion (and implications)
- Data Poisoning:
 - Indiscriminate Attacks
 - Support vector machines (SVMs)
 - Regression models
 - Conclusion (and implications)

Nelson *et al.*, Exploiting Machine Learning to Subvert Your Spam Filter
Rubinstein *et al.*, ANTIDOTE: Understanding and Defending against
Poisoning of Anomaly Detectors

Motivation

- Goals

- DDoS attack [[Link](#)]

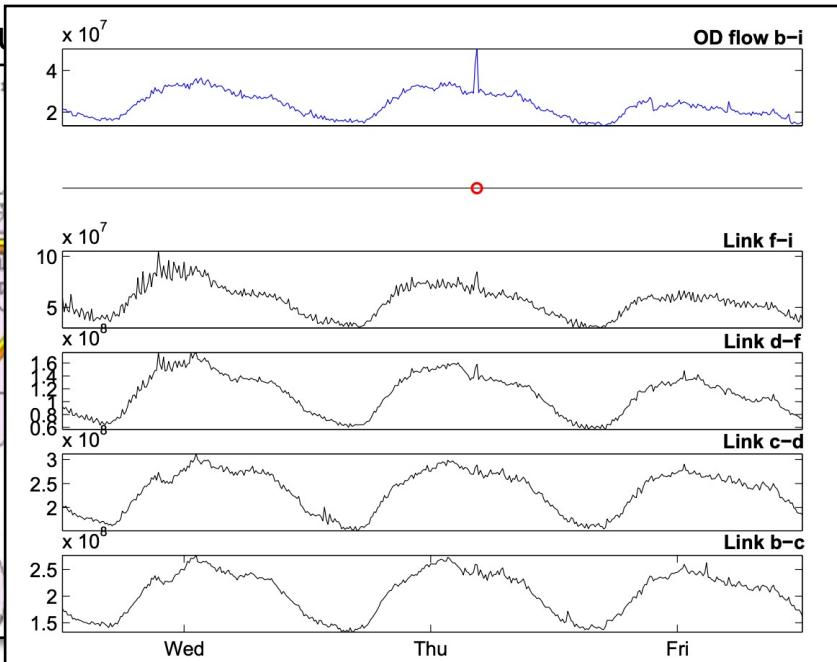
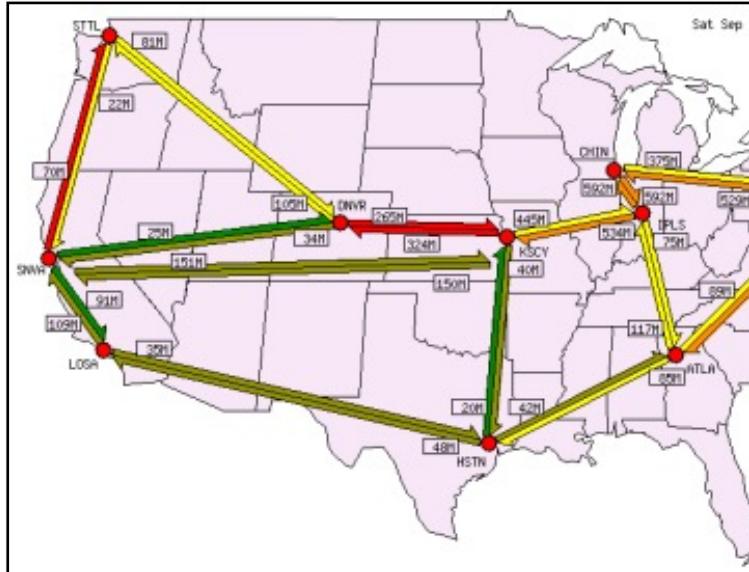


<https://edureka.co/blog/what-is-ddos-attack/>

Kang et al., Crossfire Attack, IEEE Security and Privacy 2013

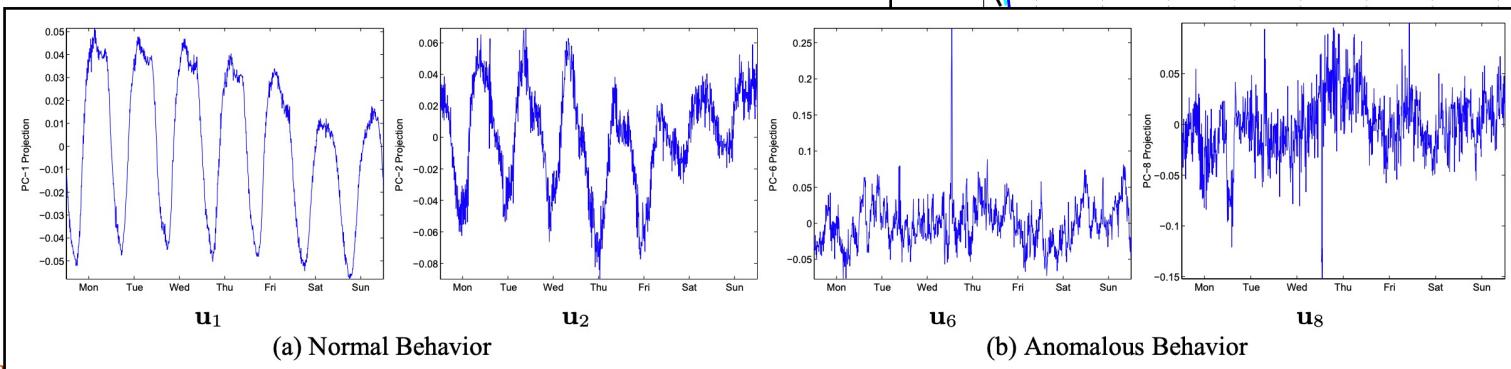
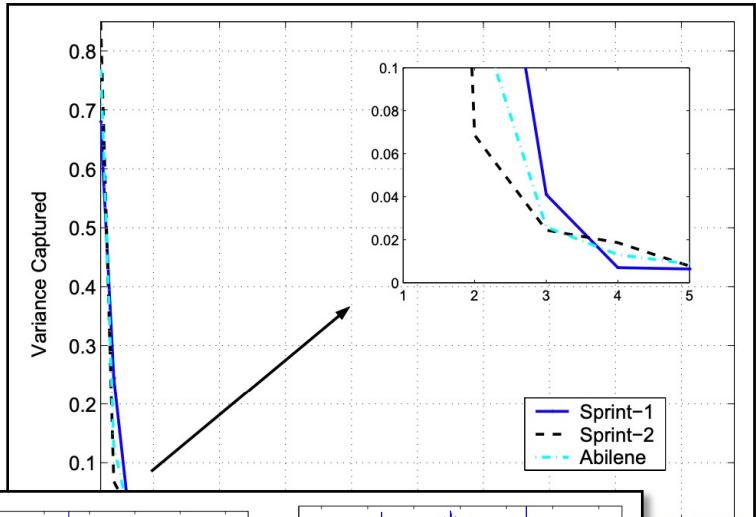
Motivation

- Goals
 - DDoS attack
 - Attacker's network traffic successfully cross an ISP's network
 - ISP Monitors in-out traffic and alert “volume”



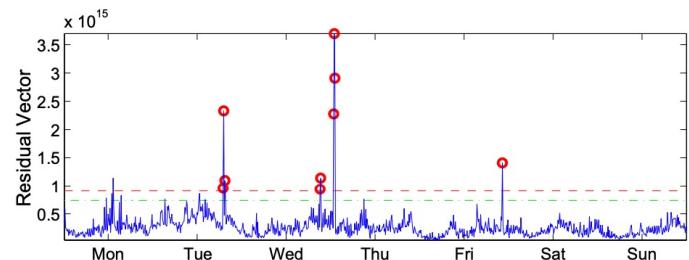
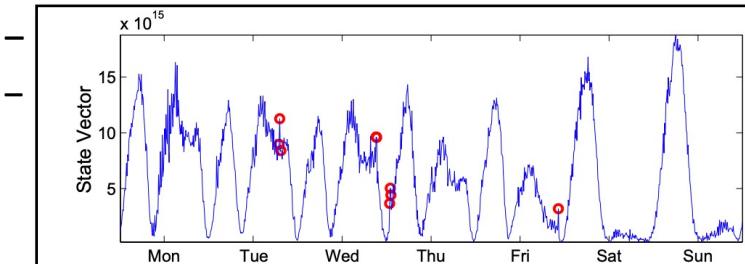
Background: PCA-based Anomaly Detector (Lakhina et al.)

- PCA (Principal Component Analysis)
 - Represent data with smaller set of variables
- PCA-based Anomaly Detection
 - $Y: T \times N$ (time series of all links)
 - Run PCA on Y
 - Find the top-K normal components
 - The rest $[N-K]$ is for detecting anomalies

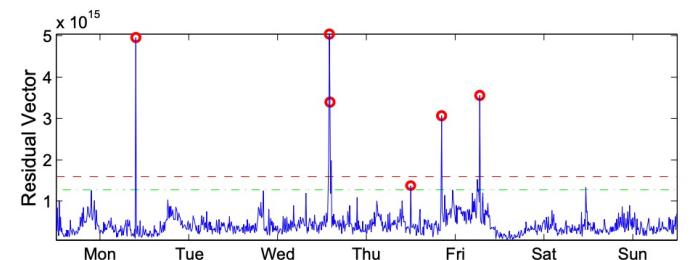
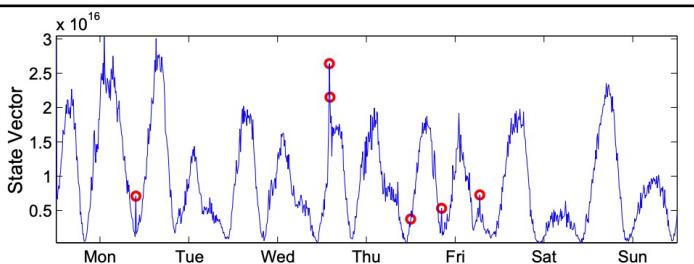


Background: PCA-based Anomaly Detector (Lakhina et al.)

- PCA (Principal Component Analysis)
 - Represent data with smaller set of variables
- PCA-based Anomaly Detection



(a) Sprint-1



(b) Sprint-2

Motivation

- Research Questions:
 - **RQ 1:** How can we **poison** the anomaly detector to launch DDoS?
 - **RQ 2:** How much this attack will be **effective**?
 - **RQ 3:** How can we **mitigate** this poisoning attacks?

Threat Model

- Goal
 - Manipulate the anomaly detector while increasing the traffic volume [~indiscriminate]
- Capability
 - Inject additional traffic (*chaff*) along the network flow
- Knowledge
 - Does not know the traffic (*uninformed* attack)
 - Know the current volume of traffic (*locally-informed* attack)
 - Know all the details about the network links (*globally-informed* attack)
- [Victim] Anomaly Detector
 - PCA retrained each week on $m - 1$ (with anomalies removed)
 - Use the trained PCA for detecting anomalies in week m

Poisoning Attack Strategies

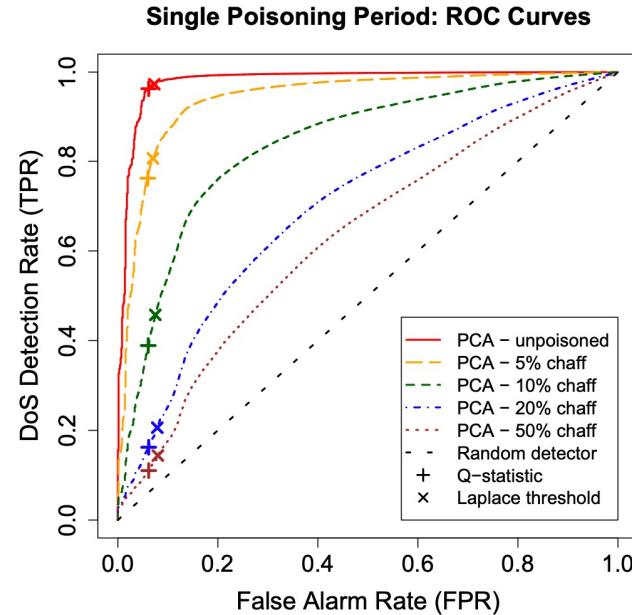
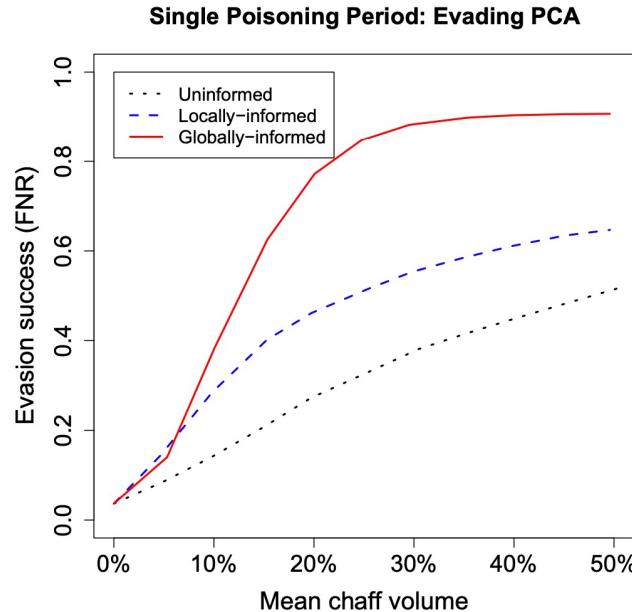
- Uninformed
 - Randomly add chaff (the amount is θ)
- Locally-informed
 - Only add chaff $(\max\{0, y_S(t) - \alpha\})^\theta$ when the traffic is already reasonably large
- Globally-informed
 - Optimize the amount of chaff
$$\begin{aligned} \max_{\mathbf{C} \in \mathbb{R}^{T \times F}} \quad & \|(\bar{\mathbf{Y}} + \mathbf{C})\mathbf{A}_f\|_2 \\ \text{s.t.} \quad & \|\mathbf{C}\|_1 \leq \theta \\ & \forall t, n \quad \mathbf{C}_{tn} \geq 0 \end{aligned}$$
- [Continuous case] Boiling Frog attack
 - Initially set the theta to a small value, and increase it over time
 - Use any of the three (informed, locally-informed, or globally-informed) to add chaff

Evaluation: Attacks

- Setup
 - Dataset: OD Flow Data from Ailene network
 - Period: Mar. 2004 – Sep. 2004 (6 months)
 - Each week: 2016 measurements x 144 networks, 5 min intervals
- Metrics
 - Detector's false negative rate (FNR)
 - Use ROC curve to show tradeoffs btw true positive rate (TPR) and FPR

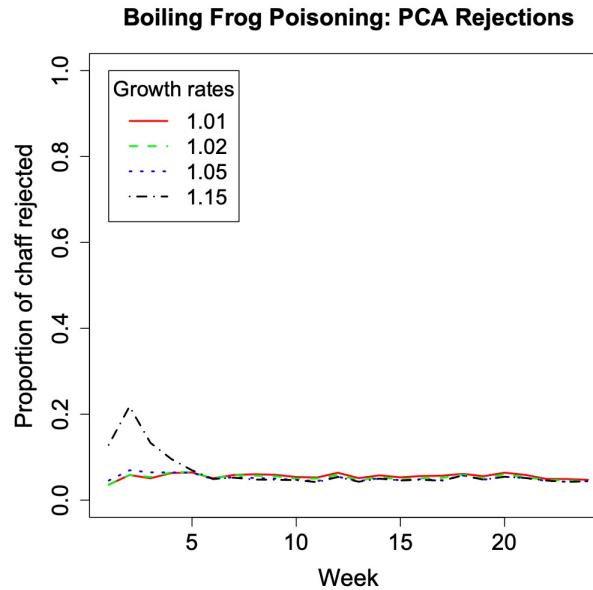
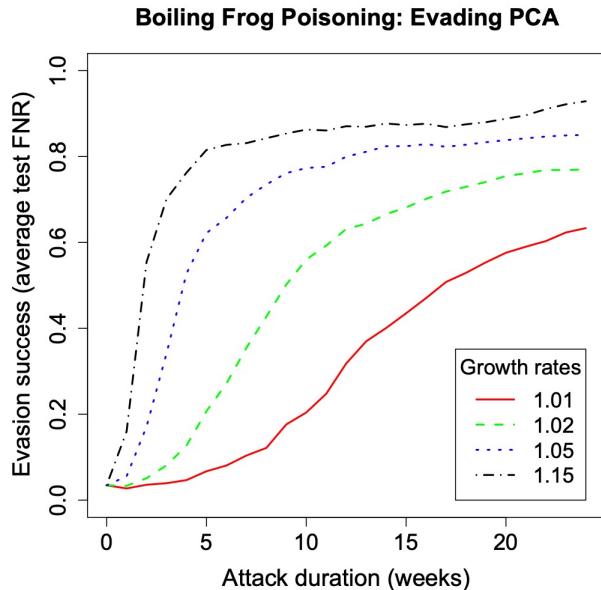
Evaluation: Attacks

- Single Poisoning Period
 - One week data for training PCA and the next one week for testing



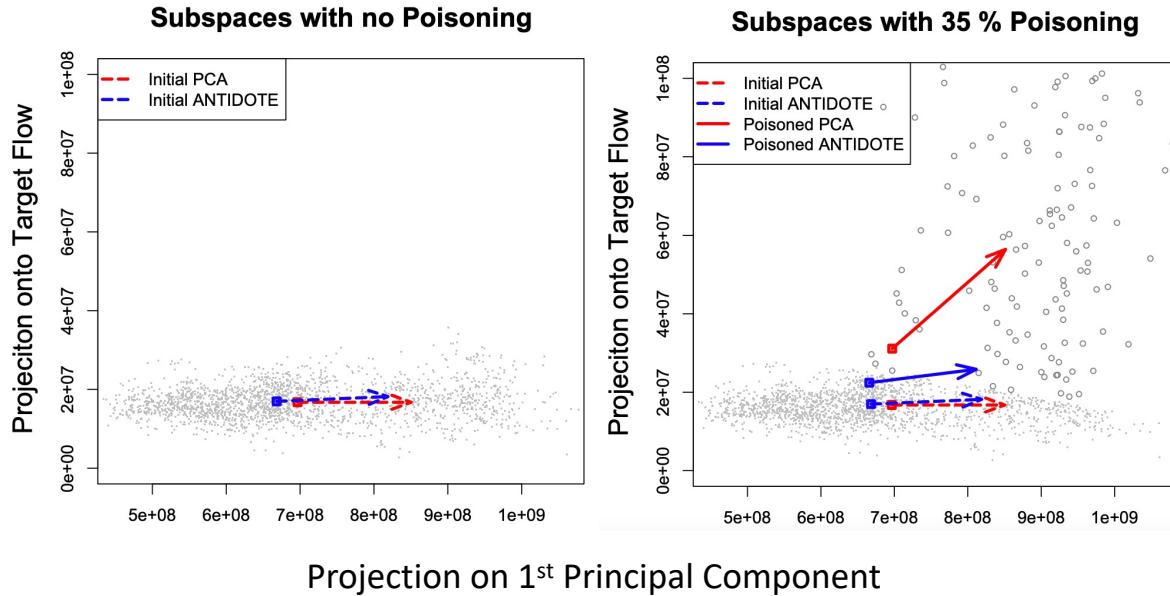
Evaluation: Attacks

- Boiling Frogs
 - Data from previous weeks for training PCA and the current week for testing



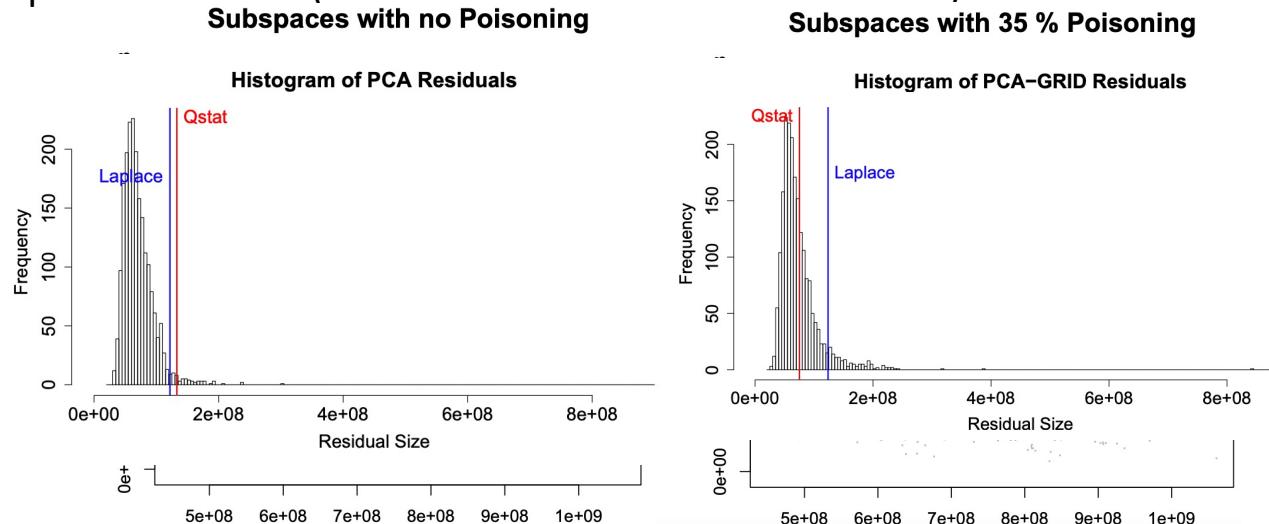
Defense: ANTIDOTE

- Robust statistics
 - Reduce the sensitivity of statistics to outliers
 - Use PCA-GRID (Croux *et al.*)



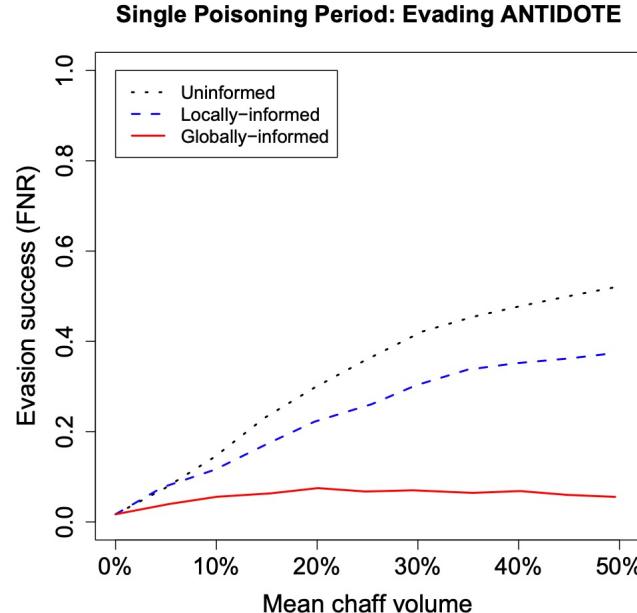
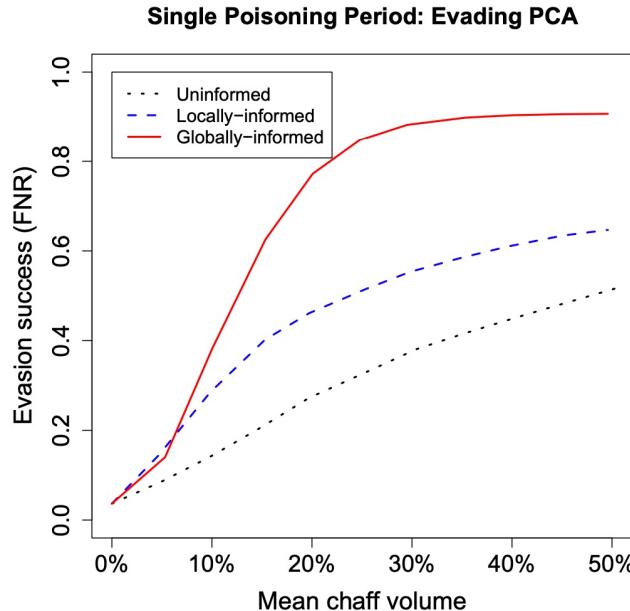
Defense: ANTIDOTE

- Robust statistics
 - Reduce the sensitivity of statistics to outliers
 - Use PCA-GRID (Croux *et al.*)
 - Use Laplace Threshold (Robust estimate for its residual threshold)



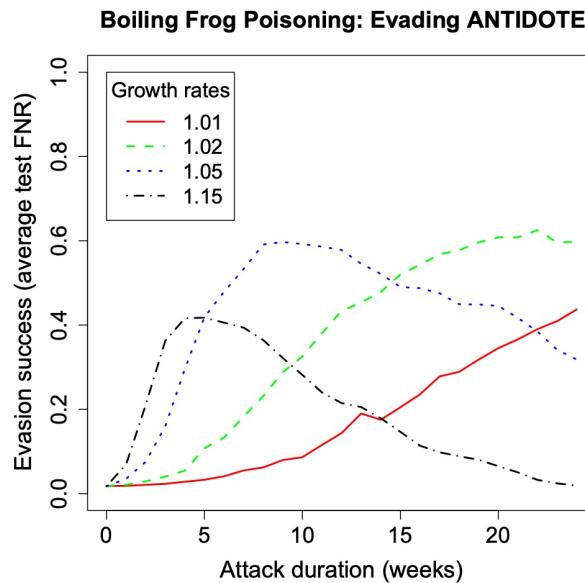
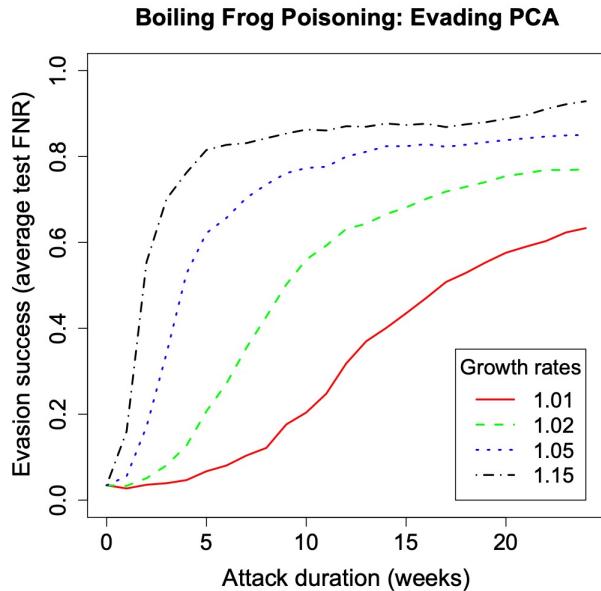
Evaluation: ANTIDOTE

- Single Poisoning Period
 - One week data for training PCA and the next one week for testing



Evaluation: Attacks

- Boiling Frogs
 - Data from previous weeks for training PCA and the current week for testing



Conclusion

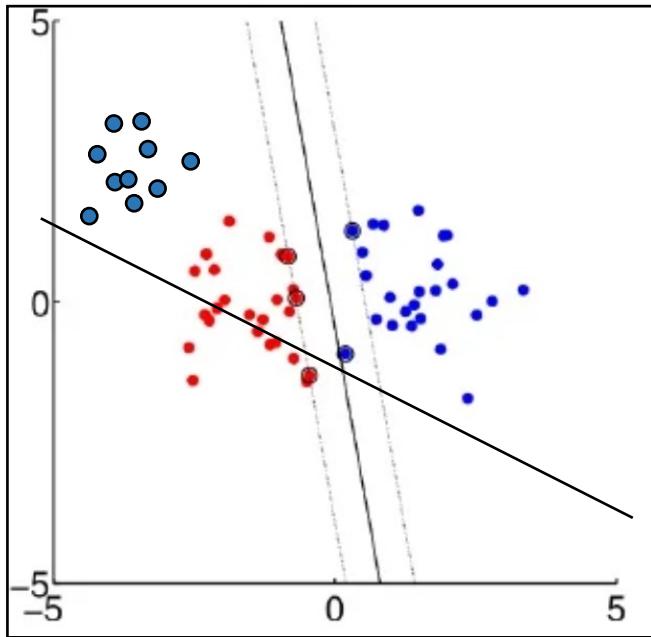
- Research Questions:
 - **RQ 1:** How can we **poison** the anomaly detector to launch DDoS?
 - Inject some additional traffic (chaff)
 - Make a detector have false estimation of normal states
 - Three-levels of knowledge: uninformed / locally-informed / globally-informed
 - Single poisoning vs. Boiling frogs
 - **RQ 2:** How much this attack will be **effective**?
 - The success increases as we increase (knowledge / % of poisons / period)
 - **RQ 3:** How can we **mitigate** this poisoning attacks?
 - ANTIDOTE: Robust statistics (PCA-GRID + Laplace threshold)

Topics for Today

- Data Poisoning
 - Exploitations
 - Spam filtering
 - DDoS detection
 - Conclusion (and implications)
- Data Poisoning
 - Indiscriminate Attacks
 - Support vector machines (SVMs)
 - Regression models
 - Conclusion (and implications)

Biggio *et al.*, Poisoning Attacks against Support Vector Machines
Jagielski *et al.*, Manipulating Machine Learning: Poisoning Attacks
and Countermeasures for Regression Learning

Revisited: Linear Models vs. DNNs



← Linear model (SVM)

Background: Support Vector Machine

- DIT [[Link](#)]
 - 1: let's put green points
 - 2: let's put red points on the other side
 - 3: let's put red points closer to the green cluster
 - 4: let's put red points in the middle of the green cluster
 - 5: let's use another kernel.

Threat Model

- Goal
 - Indiscriminate attack
 - Find a point (x_c, y_c) , whose addition to D_{tr} decreases a model's acc.
- Capability
 - Train a model f on D_{tr}
 - Inject the point (x_c, y_c) into D_{tr}
- Knowledge
 - D_{tr} : training data
 - D_{val} : validation data (where we pick the poison)
 - f : a (linear) SVM and its parameters a_i, b
 - A : training algorithm (e.g., Sub-Gradient Descent)

Proposed Attack on SVM!

Algorithm 1 Poisoning attack against SVM

Input: \mathcal{D}_{tr} , the training data; \mathcal{D}_{val} , the validation data; y_c , the class label of the attack point; $x_c^{(0)}$, the initial attack point; t , the step size.

Output: x_c , the final attack point.

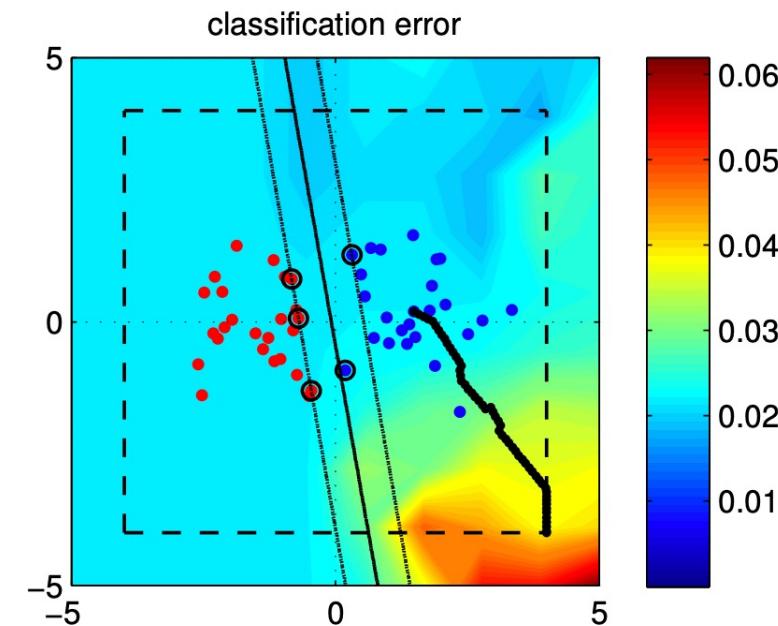
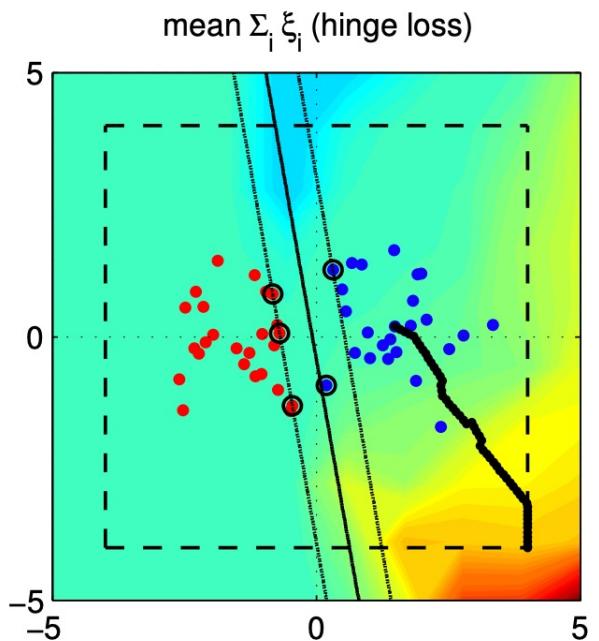
- 1: $\{\alpha_i, b\} \leftarrow$ learn an SVM on \mathcal{D}_{tr} . // train an SVM on the clean data
- 2: $k \leftarrow 0$.
- 3: **repeat**
- 4: Re-compute the SVM solution on $\mathcal{D}_{\text{tr}} \cup \{x_c^{(p)}, y_c\}$ using incremental SVM (e.g., Cauwenberghs & Poggio, 2001). This step requires $\{\alpha_i, b\}$. // train an SVM with the poison
- 5: Compute $\frac{\partial L}{\partial u}$ on \mathcal{D}_{val} according to Eq. (10). // compute the gradient
- 6: Set u to a unit vector aligned with $\frac{\partial L}{\partial u}$.
- 7: $k \leftarrow k + 1$ and $x_c^{(p)} \leftarrow x_c^{(p-1)} + tu$ // update the poison, to increase the loss
- 8: **until** $L(x_c^{(p)}) - L(x_c^{(p-1)}) < \epsilon$ // stop if the loss doesn't increase more than ϵ
- 9: **return:** $x_c = x_c^{(p)}$

Evaluation

- Setup
 - Datasets
 - Artificial data: Gaussian dist. [$N(-1.5, 0.6^2)$ vs. $N(1.5, 0.6^2)$]
 - Real data: MNIST [1 vs. 7 | 8 vs. 9 | 0 vs. 4]
 - Model(s)
 - SVM [Linear vs. RBF-Kernel]

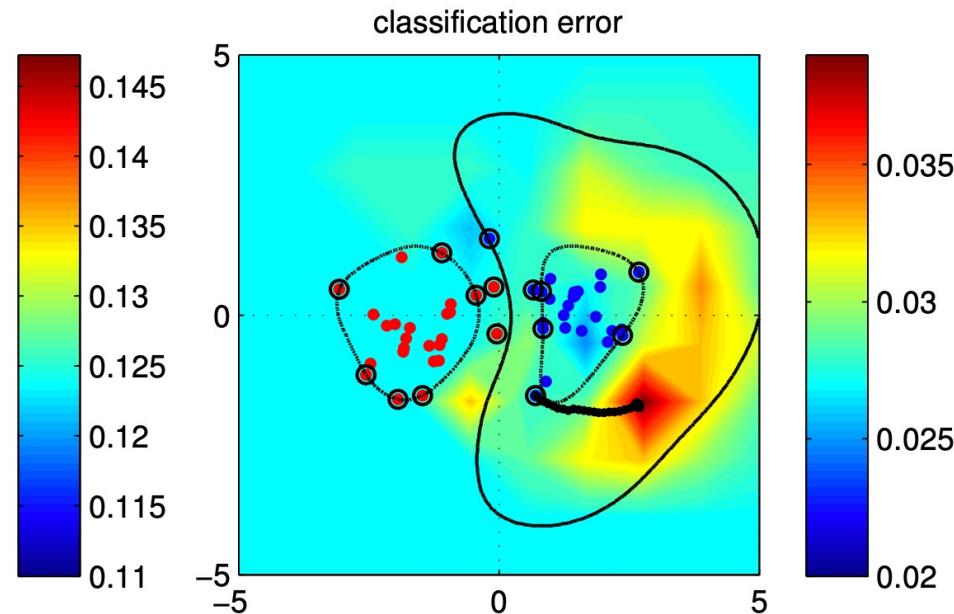
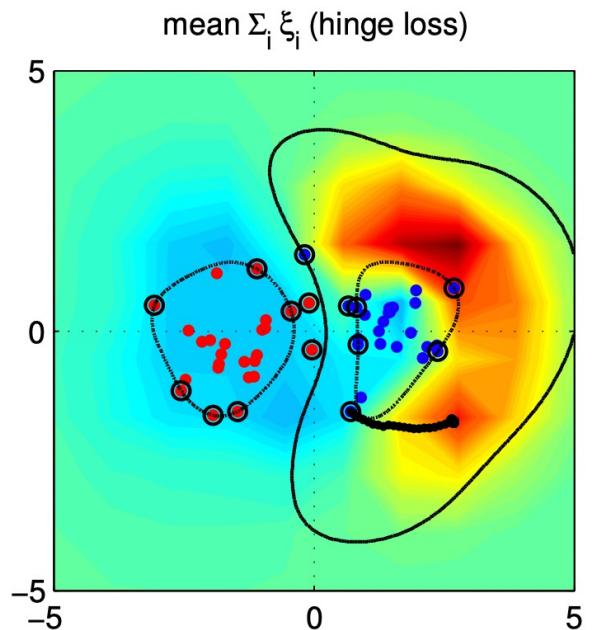
Evaluation: Artificial Data

- Linear SVM



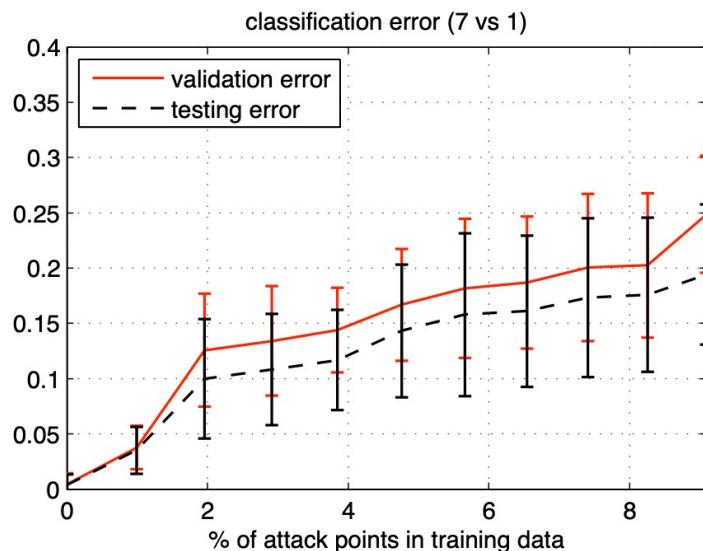
Evaluation: Artificial Data

- SVM w. RBF Kernel



Evaluation: MNIST

- Linear SVM



- Results

- Use a *single* poison
- Error increases by 15 – 20%
- Increasing # poisons leads to a higher error

Topics for Today

- Data Poisoning
 - Exploitations
 - Spam filtering
 - DDoS detection
 - Conclusion (and implications)
- Data Poisoning
 - Indiscriminate Attacks
 - Support vector machines (SVMs)
 - Regression models
 - Conclusion (and implications)

Biggio *et al.*, Poisoning Attacks against Support Vector Machines
Jagielski *et al.*, Manipulating Machine Learning: Poisoning Attacks
and Countermeasures for Regression Learning

Background: Regression Models

- Regression Models [[Demo](#)]
 - DIT
 - 1. let's add some more points
 - 2. let's see how much error ($RMSE$) it increases
 - In the Paper
 - Ordinary Least Squares (OLS)
 - Ridge regression
 - LASSO
 - Elastic-net regression

Threat Model

- Goal
 - Indiscriminate attack (increase the error on D_{val})
- Capability
 - Train a model f on D_{tr}
 - Inject p poisons into the training set ($N(D_{tr}) = n + p$)
- Knowledge [White-box vs. Black-box]
 - D_{tr} : training data (black-box adversary only has partial knowledge of D_{tr})
 - D_{val} : validation data
 - f : a model and its parameters (black-box attacker doesn't know the parameters)
 - L : training algorithm

Attack Formulation: Bi-level Optimization

$$\begin{aligned} \arg \max_{\mathcal{D}_p} \quad & \mathcal{W}(\mathcal{D}', \theta_p^*) , \\ \text{s.t.} \quad & \theta_p^* \in \arg \min_{\theta} \mathcal{L}(\mathcal{D}_{\text{tr}} \cup \mathcal{D}_p, \theta) \end{aligned}$$

- Outer-optimization: maximize the error of a model on the validation data
- Inner-optimization: minimize the model's error on the training data

Proposed Attack on Regression Models!

Algorithm 1 Poisoning Attack Algorithm

Input: $\mathcal{D} = \mathcal{D}_{\text{tr}}$ (white-box) or \mathcal{D}'_{tr} (black-box), $\mathcal{D}', \mathcal{L}, \mathcal{W}$,
the initial poisoning attack samples $\mathcal{D}_p^{(0)} = (\mathbf{x}_c, y_c)_{c=1}^p$, a
small positive constant ε .

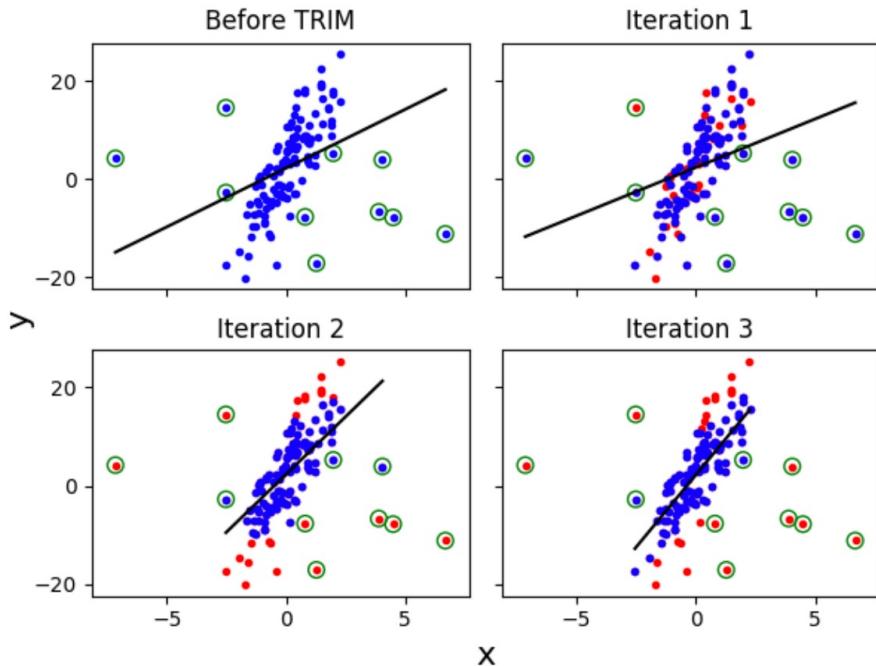
```
1:  $i \leftarrow 0$  (iteration counter)
2:  $\boldsymbol{\theta}^{(i)} \leftarrow \arg \min_{\boldsymbol{\theta}} \mathcal{L}(\mathcal{D} \cup \mathcal{D}_p^{(i)}, \boldsymbol{\theta})$  // train a model on the contaminated data
3: repeat
4:    $w^{(i)} \leftarrow \mathcal{W}(\mathcal{D}', \boldsymbol{\theta}^{(i)})$ 
5:    $\boldsymbol{\theta}^{(i+1)} \leftarrow \boldsymbol{\theta}^{(i)}$ 
6:   for  $c = 1, \dots, p$  do
7:      $\mathbf{x}_c^{(i+1)} \leftarrow \text{line\_search} \left( \mathbf{x}_c^{(i)}, \nabla_{\mathbf{x}_c} \mathcal{W}(\mathcal{D}', \boldsymbol{\theta}^{(i+1)}) \right)$  // update poisons to increase the loss of the model
8:      $\boldsymbol{\theta}^{(i+1)} \leftarrow \arg \min_{\boldsymbol{\theta}} \mathcal{L}(\mathcal{D} \cup \mathcal{D}_p^{(i+1)}, \boldsymbol{\theta})$ 
9:      $w^{(i+1)} \leftarrow \mathcal{W}(\mathcal{D}', \boldsymbol{\theta}^{(i+1)})$ 
10:     $i \leftarrow i + 1$ 
11: until  $|w^{(i)} - w^{(i-1)}| < \varepsilon$  // stop when the model doesn't change more than  $\varepsilon$ 
```

Output: the final poisoning attack samples $\mathcal{D}_p \leftarrow \mathcal{D}_p^{(i)}$

Proposed Defense: TRIM

Algorithm 2 [TRIM algorithm]

```
1: Input: Training data  $\mathcal{D} = \mathcal{D}_{\text{tr}} \cup \mathcal{D}_p$  with  $|\mathcal{D}| = N$ ;  
   number of attack points  $p = \alpha \cdot n$ .  
2: Output:  $\theta$ .  
3:  $\mathcal{I}^{(0)} \leftarrow \{1, \dots, N\}$  /* First train with all samples */  
4:  $\theta^{(0)} \leftarrow \arg \min_{\theta} \mathcal{L}(\mathcal{D}^{\mathcal{I}^{(0)}}, \theta)$  /* Initial estimation of  $\theta$ */  
5:  $i \leftarrow 0$  /* Iteration count */  
6: repeat  
7:    $i \leftarrow i + 1$ ;  
8:    $\mathcal{I}^{(i)} \leftarrow$  subset of size  $n$  that min.  $\mathcal{L}(\mathcal{D}^{\mathcal{I}^{(i)}}, \theta^{(i-1)})$   
9:    $\theta^{(i)} \leftarrow \arg \min_{\theta} \mathcal{L}(\mathcal{D}^{\mathcal{I}^{(i)}}, \theta)$  /* Current estimator */  
10:   $R^{(i)} = \mathcal{L}(\mathcal{D}^{\mathcal{I}^{(i)}}, \theta^{(i)})$  /* Current loss */  
11: until  $i > 1 \wedge R^{(i)} = R^{(i-1)}$  /* Convergence condition*/  
12: return  $\theta^{(i)}$  /* Final estimator */.
```



Evaluation

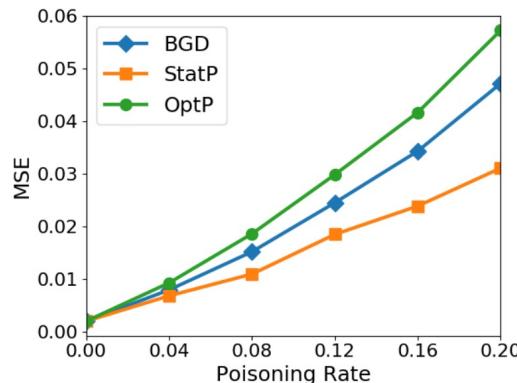
- Setup
 - Datasets: Health care | Loan | Housing
 - Models
 - Ordinary Least Square (OLS)
 - Ridge regression
 - LASSO
 - Elastic-net regression
 - Attacks
 - OptP | StatP | BGD (Prior work by Xiao *et al.*)
 - Defenses
 - Huber | RANSAC | Chen *et al.* | RONI | TRIM

Evaluation

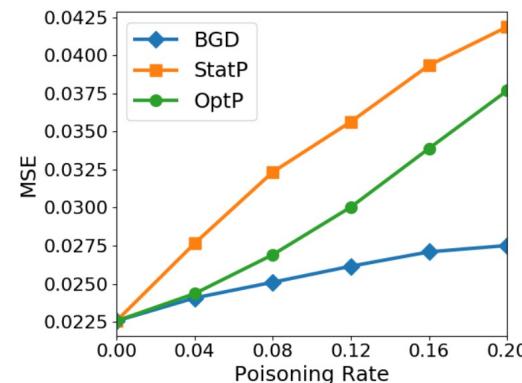
- Results Summary

- Attacks

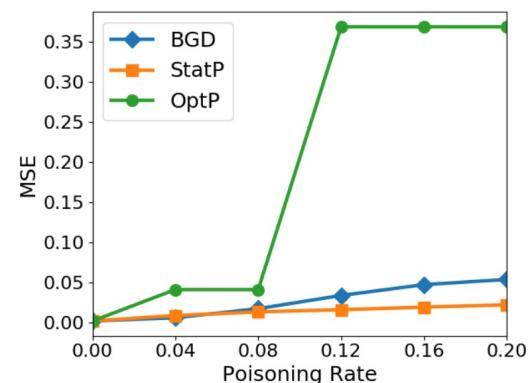
- OptP > StatP, BGD (Prior work)
 - StatP, BGD: varies from datasets
 - StatP > OptP: computational efficiency; StatP still shows a reasonable success rate
 - Poisons transfer: crafted on one model works for the three others



(a) Health Care Dataset



(b) Loan Dataset



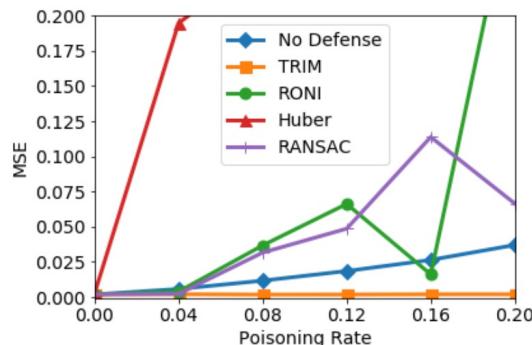
(c) House Price Dataset

Evaluation

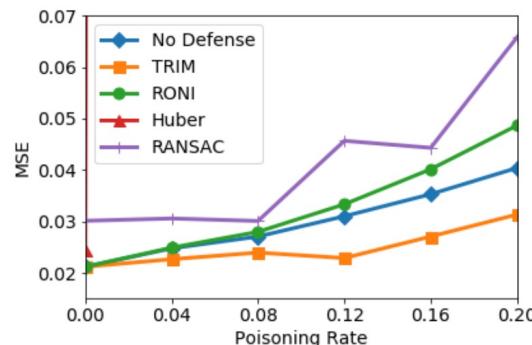
- Results Summary

- Defenses

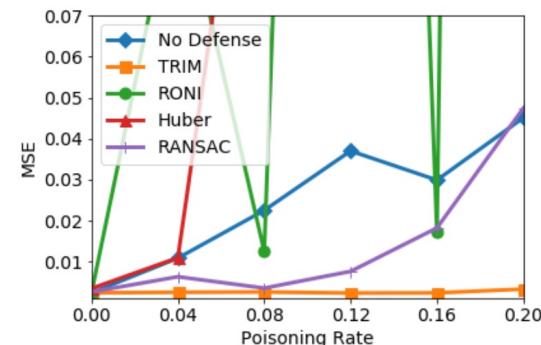
- TRIM > Huber | RANSAC | Chen *et al.* | RONI
 - TRIM is computationally efficient (< 0.02 seconds on the House dataset)
 - Prior work's defenses sometimes increase errors



(a) Health Care Dataset



(b) Loan Dataset



(c) House Price Dataset

Topics for Today

- Data Poisoning
 - Exploitations
 - Spam filtering
 - DDoS detection
 - Conclusion (and implications)
- Data Poisoning
 - Indiscriminate Attacks
 - Support vector machines (SVMs)
 - Regression models
 - [Now] Conclusion (and implications)

Thank You!

Mon/Wed 12:00 – 1:50 pm

Sanghyun Hong

<https://secure-ai.systems/courses/MLSec/W22>



Oregon State
University

SAIL
Secure AI Systems Lab